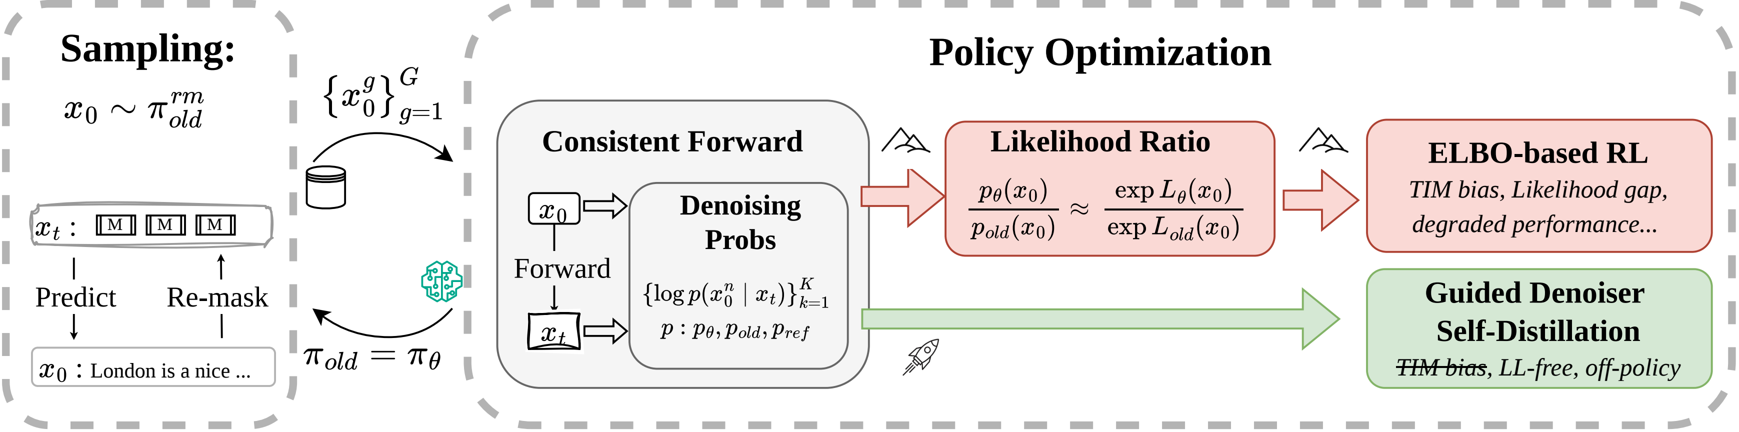

Sampling:
$x_0 \sim \pi_{old}^{rm}$
x_t : [M] [M] [M]
Predict
Re-mask
x_0 : London is a nice ...
$\{x_0^g\}_{g=1}^G$
$\pi_{old} = \pi_\theta$
Policy Optimization
Consistent Forward
x_0
Forward
x_t
Denoising Probs
$\{\log p(x_0^n \mid x_t)\}_{k=1}^K$
$p : p_\theta, p_{old}, p_{ref}$
Likelihood Ratio
$\frac{p_\theta(x_0)}{p_{old}(x_0)} \approx \frac{\exp L_\theta(x_0)}{\exp L_{old}(x_0)}$
ELBO-based RL
TIM bias, Likelihood gap, degraded performance...
Guided Denoiser Self-Distillation
~~TIM bias~~, LL-free, off-policy